

MENTICATION

Thought or Fact

Complete product, visual, clinical and engineering handoff

Purpose: enable a new product, design or engineering team to understand, rebuild and safely extend the intervention without needing prior project context.

What this intervention is

Thought or Fact is a guided cognitive appraisal intervention for an everyday upsetting thought. It separates an observation from interpretations, feelings and predictions; surfaces possible thinking shortcuts without diagnosing the person; gathers evidence for and against; and helps the person leave with a more balanced, editable perspective.

Current production flow


Readiness - Thought capture - Precise claim - Thinking-pattern suggestions - Statement type - Evidence for - Evidence against - Other explanations - What remains open - Balanced perspective plus final rating - Next step or handoff.

Core experience rules

Rule	Implementation meaning
User authorship	The original thought remains visible and the balanced thought is editable. No forced positive reframing.
Uncertainty	Patterns are suggestions, not verdicts. The app never labels a concern false.
One decision at a time	Each screen has one primary decision. Avoid duplicate ratings and redundant evidence overview screens.
Safety	Grounding and Test the Prediction routes remain available; testing requires an explicit safety confirmation.
Privacy	The working experience is local-first. Do not add claims about AI processing, storage or deletion without the corresponding release infrastructure.

Live screen catalogue

The following atlas is captured directly from the running production app at the 390 x 844 mobile viewport. It documents the 12 key live states in the streamlined experience: the redundant evidence overview is removed, statement sorting is one tap, and the final rating is integrated into the balanced-thought screen.



Screen-by-screen implementation notes

01 Thought capture

Capture the thought in the person's own language. Case-file textarea, voice trigger, privacy label, Open case and Ground first.

02 Voice capture

Permission-aware voice capture. Recording, elapsed time, cancellation and no retained-audio claim only when true.

03 Claim confirmation

Keep original words alongside an editable claim. Continue with either version.

04 Baseline rating

Original board baseline rating. In the current streamlined build, use the Reset start rating rather than duplicating this screen.

05 Charge sheet

Zero-to-many thinking patterns. Suggested cards must point to exact language and remain optional.

06-11 Evidence work

Production uses four concise sequential surfaces: evidence for, evidence against, other explanations and what remains open.

12-13 Questioning

Original configurable cross-examination design. Preserve only if a validated question bank is supplied.

14-16 Balanced perspective

Production uses one editable balanced-thought screen with the final rating integrated below it.

17-22 Exit and handoffs

Leave here, grounding, practical action and Test the Prediction. Handoffs must preserve the user's context without forcing a verdict.

Visual design system

Token	Specification
Canvas	390 px flagship mobile viewport; responsive 320-430 px; safe-area aware.
Palette	Deep teal #071B1A, elevated teal #0D302C, emerald #0D4B43, cream #F3E8D2, antique gold #B99350, oxblood #58212B.
Typography	Fraunces/Georgia editorial serif for display headings and thought content; Hanken Grotesk or sans for labels and controls. Body at least 16 px in app.
Spacing	24 px page margins, 8 px spacing base, 16-24 px component padding, 16 px minimum between unrelated controls.
Controls	Primary 56 px targets; all controls at least 44 px; visible focus; selected state must not rely only on colour.
Material language	Cream is reserved for authored content and foreground paper; brass borders are fine and restrained; avoid generic rounded-card stacks.
Motion	180-260 ms standard; only opacity/transform; reduced-motion falls back to crossfade.

Visual assets

Supplied artwork: asset-01-case-file.png, asset-02-evidence-scale.png and asset-03-completion-seal.png. UI icons are code-native Lucide vectors: back, close, microphone, lock, evidence categories, editing, plus/minus and handoff symbols. Do not substitute emojis or custom CSS illustrations.

Guided thinking-pattern suggestions

The release experience is designed for AI-assisted analysis. The UI never needs to announce whether analysis is AI or local; it should simply explain why a pattern is suggested and let the person decide. Every suggested pattern must expose the exact triggering language in the claim.

Pattern	Typical language signal	Balanced response direction
Mind reading	they think, everyone thinks, nobody likes me	Name uncertainty about other minds and ask what can actually be known.
Jumping to conclusions	they will, must mean, definitely	Turn a conclusion back into a possibility.
Overgeneralising	all people, everyone, nobody, always, never	Move from an absolute group conclusion to specific people or moments.
Emotional reasoning	I feel, I hate, I am anxious	Validate the feeling without using it as proof.
Discounting positives	does not count, does not matter	Keep exceptions and disconfirming evidence in the picture.
Personalising	my fault, because of me	Make room for multiple causes.

Current code provides transparent local phrase matching as a fallback. A release AI service should return the same shape: pattern ID, exact source spans, short reason, confidence, and a safe fallback when no pattern is appropriate.

Personalised balanced thought

The conclusion must be short, specific and credible. It should not concatenate generic CBT statements. The output should use: the original claim, selected patterns, evidence against, alternative explanations and open questions. It must be editable.

Input	Example
Original thought	I hate all people in Melbourne because I just feel like everyone is cooked.
Patterns selected	Overgeneralising; emotional reasoning.
Evidence against	My sister and two friends have treated me well.
Balanced thought	I am feeling fed up with people in Melbourne right now. That feeling is real, but it does not prove every person is the same. My sister and two friends have treated me well is an exception worth holding alongside it.

AI integration contract

Input to the AI service should be an ephemeral case object with raw thought, active claim, selected patterns, evidence entries and user-editable final thought. Output should include only allowed pattern IDs, exact quotation spans, a brief rationale and a maximum three-sentence balanced thought. Validate output against the approved taxonomy, preserve uncertainty, block diagnoses and crisis advice, and always permit manual editing or skipping.

Engineering map

File	Responsibility
src/pages/ResetFlow.jsx	Entry readiness, Thought capture folder, voice permission state, initial Reset rating handoff and intervention routing.
src/components/ThoughtOrFactExperience.jsx	The staged intervention: claim confirmation, patterns, statement type, evidence lanes, balanced thought, final rating and handoffs.
src/lib/thoughtOrFactState.js	Draft normalisation, non-identifying learning record, language signal matching and balanced-thought builder.
src/styles/thought-or-fact.css	Dedicated visual system, safe areas, responsive layout, reduced motion and high-contrast states.
src/lib/eliteInterventionSystem.test.js	Focused regression suite for routing, safety, privacy language, state normalisation and cognitive pattern behaviour.

State and persistence

Keep raw thought, refined claim, thinking patterns, statement classification, evidence arrays, balanced thought, rating and selected handoff separate. Normalise drafts before local persistence. Recommendation learning records must retain only aggregate metadata such as classification and rating shift - not the person's raw words or evidence.

Accessibility and device behaviour

Use semantic buttons, labels, details/summary disclosures and range inputs. Maintain 16 px minimum editable input text to prevent iOS zoom. Support keyboard focus, screen-reader announcements, safe-area insets, reduced motion, high contrast and large text. A small mobile screen must not require scrolling merely to reveal a required choice.

Release acceptance checklist

Area	Must pass before release
Visual QA	Inspect each state at 390x844 plus 320, 375, 430, tablet and desktop. Check contrast, alignment, touch targets and no horizontal overflow.
Flow QA	One start rating and one end rating only. One-tap statement categorisation. No redundant evidence overview. All handoffs preserve state.
AI QA	Every suggestion includes exact source spans; output stays tentative, editable and situation-specific; no diagnosis or invented facts.
Privacy QA	Claims about deletion, AI processing, retention or storage must match the shipped architecture and policy.
Safety QA	Grounding, leave and Test the Prediction boundaries work; unsafe test scenarios never receive an experiment prompt.
Engineering QA	Run focused tests, lint, production build and native sync when preparing Capacitor builds.

How to rebuild from this handoff

Start with the visual system and the first mobile screen. Implement one stage at a time, write a focused behaviour test first, inspect the rendered screen at the target viewport, then proceed. Keep the component and CSS scoped to Thought or Fact. Integrate the AI service behind the documented data contract only after retention, safety and policy decisions are final.

Source material: Mentication Thought or Fact build specification and supplied screen mock-ups. This document records the currently implemented streamlined production flow where it intentionally differs from the original 22-screen board.